

Ubuntu 20.04 LTS

1) Create 3 ubuntu VM, 1 control plane (master) and 2 worker nodes.

2) Ports enable for master node: (Inbound)

TCP 6443 → For Kubernetes API server

TCP 2379–2380 → For etcd server client API

TCP 10250 → For Kubelet API

TCP 10259 → For kube-scheduler

TCP 10257 → For kube-controller-manager

TCP 22 → For remote access with ssh

UDP 8472 → Cluster-Wide Network Comm. — Flannel VXLAN

3) Ports enabled for worker nodes (Inbound)

TCP 10250 → For Kubelet API

TCP 30000–32767 → NodePort Services†

TCP 22 → For remote access with ssh

UDP 8472 → Cluster-Wide Network Comm. — Flannel VXLAN

4) Login into the controller instance and change its hostname for our convenience

```
sudo hostnamectl set-hostname k8s-controller
```

Now logout and log back in to see the change.

Similarly, do this for both the worker nodes.

```
sudo hostnamectl set-hostname k8s-worker-1
```

```
sudo hostnamectl set-hostname k8s-worker-2
```

5) Now, we will declare the known hosts in all the three nodes so that the name that we give to each host could be mapped to their **private** IP. It is done so that we dont have to be confused and memorize the ip address.

We use private IP because incase of EC2, everytime the machine restarts the public IP changes.

```
sudo nano /etc/hosts
```

We need to enter the below line in each node having all the hosts as a key value pair (where key is private ip and value is the name given by us (any random name))

Example:

```
172.31.80.250 k8s-controller
```

```
172.31.82.32 k8s-worker-1
```

```
172.31.92.113 k8s-worker-2
```

Once we enter the hosts we need to logout and log back in.

6) Run the below commands on all the nodes (both controller and worker)

On all nodes, set up Docker Engine and containerd. You will need to load some kernel modules and modify some system settings as part of this process

```
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
overlay
br_netfilter
EOF
```

```
sudo modprobe overlay
sudo modprobe br_netfilter
```

sysctl params required by setup, params persist across reboots

```
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
EOF
```

Apply sysctl params without reboot

```
sudo sysctl --system
```

Set up the Docker Engine repository

```
sudo apt-get update && sudo apt-get install -y ca-certificates curl gnupg lsb-release
apt-transport-https
```

Add Docker's official GPG key

```
sudo mkdir -m 0755 -p /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o
/etc/apt/keyrings/docker.gpg
```

Set up the repository

```
echo \
"deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg]
https://download.docker.com/linux/ubuntu \
$(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

Update the apt package index

```
sudo apt-get update
```

Install Docker Engine, containerd, and Docker Compose

```
VERSION_STRING=5:23.0.1-1~ubuntu.20.04~focal  
sudo apt-get install -y docker-ce=$VERSION_STRING docker-ce-cli=$VERSION_STRING  
containerd.io docker-buildx-plugin docker-compose-plugin
```

Add your 'cloud_user' to the docker group

```
sudo usermod -aG docker <USER>
```

Example:

```
sudo usermod -aG docker ubuntu
```

Log out and log back in so that your group membership is re-evaluated

Make sure that 'disabled_plugins' is commented out in your config.toml file

```
sudo sed -i 's/disabled_plugins/#disabled_plugins/' /etc/containerd/config.toml
```

Restart containerd

```
sudo systemctl restart containerd
```

On all nodes, disable swap.

```
sudo swapoff -a
```

On all nodes, install kubeadm, kubelet, and kubectl

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -
```

```
cat << EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list  
deb https://apt.kubernetes.io/ kubernetes-xenial main  
EOF
```

```
sudo apt-get update && sudo apt-get install -y kubelet=1.24.0-00 kubeadm=1.24.0-00  
kubectl=1.24.0-00
```

```
sudo apt-mark hold kubelet kubeadm kubectl
```

7) Now go back to the controller plane (Master), and run the below command

On the control plane node only, initialize the cluster and set up kubectl access

```
sudo kubeadm init --pod-network-cidr 192.168.0.0/16 --kubernetes-version 1.24.0
```

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Verify the cluster is working

```
kubectl get nodes
```

Install the Calico network add-on

```
kubectl apply -f
```

<https://raw.githubusercontent.com/projectcalico/calico/v3.25.0/manifests/calico.yaml>

Get the join command (this command is also printed during kubeadm init . Feel free to simply copy it from there)

```
kubeadm token create --print-join-command
```

The above command prints a command for the worker nodes to join the cluster. Paste the returned command and execute it in all the worker nodes.

8) Copy the command returned from the last command in the controller and paste it into all the worker nodes

NOTE: you might need to use 'sudo' if it denies and gives error

9) Go back to controller node and check the nodes in the cluster

```
kubectl get nodes
```

NOTE: It might take a couple of minutes to reflect and come in ready mode.

Docs on installing Kubeadm:

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

To see more on creating a cluster using kubeadm:

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/>

ORIGINAL content:

Building a Kubernetes Cluster

Lesson URL:

<https://learn.acloud.guru/course/introduction-to-kubernetes/learn/9c48bcf2-2573-485f-8906-6977bed23fc0/10592c01-d22c-4620-bace-28de4de4cf4b/watch>

Relevant Documentation

- Installing kubeadm:

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

- Creating a cluster with kubeadm:

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/>

Lesson Reference

If you are using cloud playground, create three servers with the following settings:

- Distribution: Ubuntu 20.04 Focal Fossa LTS

- Size: medium

If you wish, you can set an appropriate hostname for each node.

On the control plane node:

```
sudo hostnamectl set-hostname k8s-control
```

On the first worker node:

```
sudo hostnamectl set-hostname k8s-worker1
```

On the second worker node:

```
sudo hostnamectl set-hostname k8s-worker2
```

On all nodes, set up the hosts file to enable all the nodes to reach each other using these hostnames

```
sudo vi /etc/hosts
```

On all nodes, add the following at the end of the file. You will need to supply the actual private IP address for each node

```
<control plane node private IP> k8s-control
```

```
<worker node 1 private IP> k8s-worker1
```

```
<worker node 2 private IP> k8s-worker2
```

Log out of all three servers and log back in to see these changes take effect

On all nodes, set up Docker Engine and containerd. You will need to load some kernel modules and modify some system settings as part of this process

```
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
overlay
br_netfilter
EOF
```

```
sudo modprobe overlay
sudo modprobe br_netfilter
```

sysctl params required by setup, params persist across reboots

```
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
EOF
```

Apply sysctl params without reboot

```
sudo sysctl --system
```

Set up the Docker Engine repository

```
sudo apt-get update && sudo apt-get install -y ca-certificates curl gnupg lsb-release
apt-transport-https
```

Add Docker's official GPG key

```
sudo mkdir -m 0755 -p /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o
/etc/apt/keyrings/docker.gpg
```

Set up the repository

```
echo \
"deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg]
https://download.docker.com/linux/ubuntu \
$(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

Update the apt package index

```
sudo apt-get update
```

Install Docker Engine, containerd, and Docker Compose

```
VERSION_STRING=5:23.0.1-1~ubuntu.20.04~focal
sudo apt-get install -y docker-ce=$VERSION_STRING docker-ce-cli=$VERSION_STRING
containerd.io docker-buildx-plugin docker-compose-plugin
```

Add your 'cloud_user' to the docker group

```
sudo usermod -aG docker $USER
```

Log out and log back in so that your group membership is re-evaluated

Make sure that 'disabled_plugins' is commented out in your config.toml file

```
sudo sed -i 's/disabled_plugins/#disabled_plugins/' /etc/containerd/config.toml
```

Restart containerd

```
sudo systemctl restart containerd
```

On all nodes, disable swap.

```
sudo swapoff -a
```

On all nodes, install kubeadm, kubelet, and kubectl

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -
```

```
cat << EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb https://apt.kubernetes.io/ kubernetes-xenial main
EOF
```

```
sudo apt-get update && sudo apt-get install -y kubelet=1.24.0-00 kubeadm=1.24.0-00
kubectl=1.24.0-00
```

```
sudo apt-mark hold kubelet kubeadm kubectl
```

On the control plane node only, initialize the cluster and set up kubectl access

```
sudo kubeadm init --pod-network-cidr 192.168.0.0/16 --kubernetes-version 1.24.0
```

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Verify the cluster is working

```
kubectl get nodes
```

```
# Install the Calico network add-on
```

```
kubectl apply -f
```

```
https://raw.githubusercontent.com/projectcalico/calico/v3.25.0/manifests/calico.yaml
```

```
# Get the join command (this command is also printed during kubeadm init . Feel free to simply copy it from there)
```

```
kubeadm token create --print-join-command
```

```
# Copy the join command from the control plane node. Run it on each worker node as root (i.e. with sudo )
```

```
sudo kubeadm join ...
```

```
# On the control plane node, verify all nodes in your cluster are ready. Note that it may take a few moments for all of the nodes to enter the READY state
```

```
kubectl get nodes
```